



Where does each window-trained transformer generalize?
8,000 trajectories per model; color compression prevents failures from dominating


Width 1 · 4 training windows





Width 0.5 · 7 training windows



Width 0.25 · 14 training windows



Width 0.125 · 28 training windows



Typical extrapolation penalty (median and interquartile range)

