

Implied return map at $R = 1$: does the model track the true map, or the one it was trained on?
 Orange on green = it identified the map from context; orange on dashed navy = it is applying a memorized map

