



continuous input



token input



Curves below the dashed line generalize beyond the training distribution. Continuous input does; token input does not.

\*the best predictor using only maps seen in training — the analogue of the optimal in-task-distribution Bayes solution.