





Trained on the last position



Trained on every position



Square loss against the exact next state, 160k steps, 32000 trajectories, context 50 — identical either side; only the training loss differs. Evaluation reads the final position in both. Right-hand points are the five seeds, line is their median.