




Models trained at context 50 (square loss, 64 input bins, 320k steps), evaluated with attention masked to the last L positions. Error barely moves from $L = 2$ to $L = 30$, then falls two orders of magnitude over the final twenty — the model needs nearly its whole trained window to engage.