




Square loss against the exact next state, 64 input bins, 320k steps, 32000 trajectories, one model per context. $\text{traj_len} = \text{context_len} + 100$ throughout, so every arm keeps 100 windows per trajectory and the same pool.