




Square loss, 64 input bins, 64-bin output, 320k steps, 32000 trajectories, one model trained per context. traj_len = context_len + 100 throughout, so every arm keeps 99 windows per trajectory and the same pool.