




Continuous input, square loss, 320k steps, 32000 trajectories. Uniform noise of half-width  $1/(2n)$  is exactly the error distribution of quantising to  $n$  bins. Noise is baked into the dataset, not resampled, so a given occurrence is always corrupted the same way.