




Square loss against the exact next state, 64 input bins, 320k steps, 32000 trajectories, one model per context length. traj_len = context_len + 100 throughout, so every arm keeps 100 windows per trajectory and the same 3.2M pool.