



Cross-entropy loss



Square loss



Output held identical across arms and total trajectories at 32000 throughout, so only the input resolution changes. Fixed-step checkpoint. Dotted: what a model that learnt nothing achieves. Continuous input is excluded -- it is a different inductive bias, not the fine end of this ladder.