



Where does each window-trained transformer generalize?
 8,000 trajectories per model; log mean squared error colour scale; white boxes are training ranges

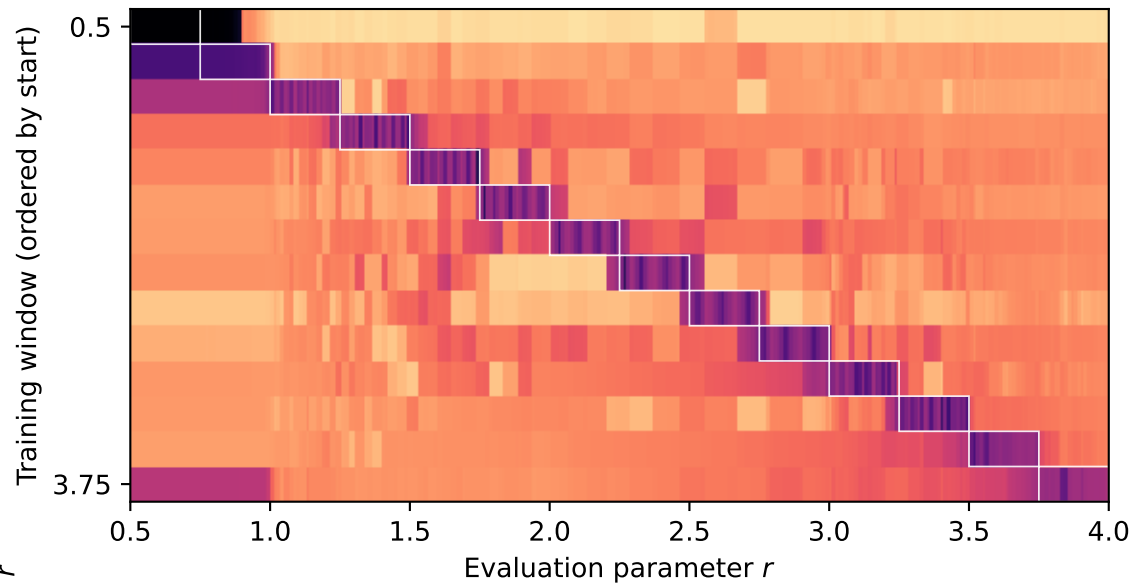
Width 1 · 4 training windows



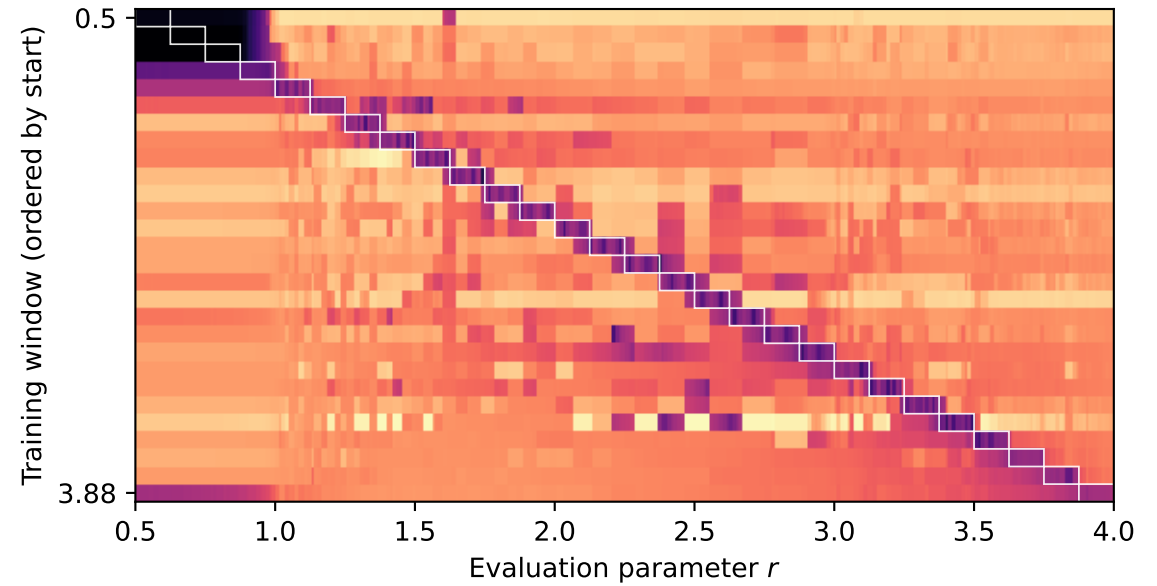
Width 0.5 · 7 training windows



Width 0.25 · 14 training windows



Width 0.125 · 28 training windows



log₁₀ (MSE); white boxes are training ranges

Typical extrapolation penalty (median and interquartile range)

